

https://drive.google.com/file/d/1v6ouvJyuzhZ80-HkZ96EoUjkt6xEUA2G/view?usp=drive_link

Pima Indians Diabetes Dataset.csv

```
import pandas as pd
import numpy as np
from google.colab import drive
drive.mount('/content/drive')
```

```
df = pd.read_csv("/content/drive/MyDrive/AI Lab(SE334)/Pima Indians Diabetes Dataset.csv")
```

Double-click (or enter) to edit

Column	Type	Status
Pregnancies	int	correct
Glucose		
BloodPressure		
SkinThickness		
Insulin		
BMI	int	correct
	int	correct
	int	correct
	int	correct
DiabetesPedigreeFunction	float	correct
Age		
Outcome		
	float	correct
	int	correct
	int	correct (binary 0/1)

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/")
```

```
#Drop rows with any missing values
df = df.dropna()
```

```
df = df.drop_duplicates()  
df = df.convert_dtypes()  
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype   ---  ----- 
 0   Pregnancies      768 non-null    Int64  
 1   Glucose          768 non-null    Int64  
 2   BloodPressure    768 non-null    Int64  
 3   SkinThickness    768 non-null    Int64  
 4   Insulin          768 non-null    Int64  
 5   BMI              768 non-null    Float64 
 6   DiabetesPedigreeFunction 768 non-null    Float64 
 7   Age              768 non-null    Int64  
 8   Outcome          768 non-null    Int64  
dtypes: Float64(2), Int64(7)
memory usage: 60.9 KB
None
```

```
print(df.columns)
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
df.isnull().sum()
```

```
0
Pregnancies
Glucose
BloodPressure
SkinThickness
Insulin
BMI
0
0
0
0
0
0
0
0
DiabetesPedigreeFunction 0
Age
0
Outcome
dtype: int64
0
```

Start coding or generate with AI.

No encoding as all data are numerical

```
#featurescalingwithstandarscale
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```
X = df.drop('Outcome', axis=1)
y = df['Outcome']
```

Feature scaling

```
scaled_x = scaler.fit_transform(X)
scaled_x
```

```
array([[ 0.63994726,  0.84832379,  0.14964075, ...,  0.20401277,
0.46849198,  1.4259954 ],
[-0.84488505, -1.12339636, -0.16054575, ..., -0.68442195,-0.36506078, -0.19067191],
[ 1.23388019,  1.94372388, -0.26394125, ..., -1.10325546,
0.60439732, -0.10558415],
...,
[ 0.3429808 ,  0.00330087,  0.14964075, ..., -0.73518964,-0.68519336, -0.27575966],
[-0.84488505,  0.1597866 , -0.47073225, ..., -0.24020459,-0.37110101,  1.17073215],
[-0.84488505, -0.8730192 ,  0.04624525, ..., -0.20212881,-0.47378505, -0.87137393]])
```

Start coding or [generate](#) with AI.

Train_test split

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(scaled_x, y, test_size=0.2, random_state
```

Logistic Regression

```
from sklearn.linear_model import LogisticRegression
LR=LogisticRegression()
```

```
LR.fit(X_train,y_train)
```

```
LogisticRegression
LogisticRegression()
```

Model Analysis of LR

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
y_pred = LR.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,4))
sns.heatmap(cm,annot=True)
```

```
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
#precision,f1,recall
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
	0.0	0.81	0.80	0.81
	1.0	0.65	0.67	0.66
accuracy				0.75
macro avg		0.73	0.74	0.73
weighted avg		0.76	0.75	0.75

```
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)
plt.figure()
plt.plot(fpr, tpr, color='darkorange')
```

```
from sklearn.model_selection import learning_curve
train_sizes, train_scores, test_scores = learning_curve(LR, scaled_x, y, cv=5, scoring='accuracy')
plt.plot(train_sizes, np.mean(train_scores, axis=1), label='Training score')
plt.plot(train_sizes, np.mean(test_scores, axis=1), label='Cross-validation score')
plt.title('Learning Curve')
plt.xlabel('Training examples')
plt.ylabel('Score')
plt.legend()
plt.show()
```

```
#Histograms (feature distribution)
df.hist(figsize=(10, 10))
plt.show()
```

```
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, color='blue')
plt.xlabel('Actual Target Values')
plt.ylabel('Predicted Target Values')
plt.title('Actual vs Predicted (Logistic Regression)')
plt.show()
```

Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
DT=DecisionTreeClassifier()
DT.fit(X_train,y_train)
```

- ▼ DecisionTreeClassifier
- DecisionTreeClassifier()

Start coding or generate with AI.

Model Analysis of DT

```
#confusion_matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
y_pred = DT.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
#precision,f1,recall
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0.0	0.83	0.78	0.80	99
1.0	0.64	0.71	0.67	55
accuracy			0.75	154
macro avg	0.73	0.74	0.74	154
weighted avg	0.76	0.75	0.76	154

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
AILAB_1111.ipynb - Colab
y_prob = DT.predict_proba(X_test)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
plt.figure()
plt.plot(fpr, tpr, color='darkorange', label=f"AUC = {roc_auc:.2f}")
plt.plot([0,1], [0,1], linestyle='--') # baseline
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Decision Tree ROC Curve")
plt.legend()
plt.show()
```

```
from sklearn.model_selection import learning_curve
train_sizes, train_scores, test_scores = learning_curve(DT, scaled_x, y, cv=5, scoring='accuracy')
plt.plot(train_sizes, np.mean(train_scores, axis=1), label='Training score')
plt.plot(train_sizes, np.mean(test_scores, axis=1), label='Cross-validation score')
plt.title('Learning Curve')
plt.xlabel('Training examples')
plt.ylabel('Score')
plt.legend()
plt.show()
```

Start coding or generate with AI.

KNN

```
from sklearn.neighbors import KNeighborsClassifier
KNN = KNeighborsClassifier(n_neighbors=5)
KNN.fit(X_train, y_train)
```

```
• KNeighborsClassifier
KNeighborsClassifier()
i
```

```
#confusion_matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
y_pred = KNN.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0.0	0.74	0.80	0.77	99
1.0	0.57	0.49	0.53	55
accuracy			0.69	154
macro avg	0.66	0.64	0.65	154
weighted avg	0.68	0.69	0.68	154

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
# Get probability for class 1 (required for ROC)
y_prob = KNN.predict_proba(X_test)[:, 1]
# ROC calculation
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
# Plot ROC Curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', label=f"AUC = {roc_auc:.2f}")
plt.plot([0,1], [0,1], linestyle='--') # baseline
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("KNN ROC Curve")
plt.legend()
plt.show()
```

```

from sklearn.model_selection import learning_curve
import matplotlib.pyplot as plt
train_sizes, train_scores, test_scores = learning_curve(
KNN,
# <-- KNN model
scaled_x,
y,
# <-- your features
# <-- your labels
cv=5,
scoring='accuracy',
train_sizes=np.linspace(0.1, 1.0, 10)
)
plt.plot(train_sizes, np.mean(train_scores, axis=1), label='Training score')
plt.plot(train_sizes, np.mean(test_scores, axis=1), label='Cross-validation score')
plt.title('KNN Learning Curve')

```

```

plt.xlabel('Training examples')
plt.ylabel('Score')
plt.legend()
plt.show()

```

```

from sklearn.ensemble import RandomForestClassifier
RF = RandomForestClassifier(n_estimators=100, random_state=42)
RF.fit(X_train, y_train)
▼
RandomForestClassifier
i
?
RandomForestClassifier(random_state=42)
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
# Predict using Random Forest
y_pred = RF.predict(X_test)
# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
# Plot
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Random Forest Confusion Matrix")
plt.show()

```

```

from sklearn.metrics import classification_report
# Make sure y_pred is predicted from Random Forest
y_pred = RF.predict(X_test)
# Print precision, recall, f1-score
print(classification_report(y_test, y_pred))
      precision    recall  f1-score   support
0.0       0.79     0.79     0.79      99
1.0       0.62     0.62     0.62      55
accuracy                           0.73     154
macro avg       0.70     0.70     0.70     154
weighted avg     0.73     0.73     0.73     154

```

```
from sklearn.model_selection import learning_curve
import matplotlib.pyplot as plt
import numpy as np
train_sizes, train_scores, test_scores = learning_curve(
    RF,
    # <-- Random Forest model
    scaled_x,
    y,
    # <-- your features
    # <-- your labels
    cv=5,
    scoring='accuracy',
    train_sizes=np.linspace(0.1, 1.0, 10)
)
plt.plot(train_sizes, np.mean(train_scores, axis=1), label='Training score')
plt.plot(train_sizes, np.mean(test_scores, axis=1), label='Cross-validation score')
plt.title('Random Forest Learning Curve')
plt.xlabel('Training examples')
plt.ylabel('Score')
plt.legend()
plt.show()
```

```
from sklearn.svm import SVC
▼
SVM_model = SVC(kernel='rbf', probability=True)    # rbf = default kernel
SVM_model.fit(X_train, y_train)
```

```
SVC
i
?
SVC(probability=True)
```

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
# Predict using SVM
y_pred = SVM_model.predict(X_test)
# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)
# Plot
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("SVM Confusion Matrix")
plt.show()
```

```
from sklearn.metrics import classification_report
# Predict using SVM
y_pred = SVM_model.predict(X_test)
# Print precision, recall, f1-score
print(classification_report(y_test, y_pred))
      precision    recall  f1-score   support
0.0       0.77     0.82     0.79      99
1.0       0.63     0.56     0.60      55
accuracy                           0.73     154
```

macro avg	0.70	0.69	0.70	154
weighted avg	0.72	0.73	0.72	154

If you run all this code like confusion matrix and ROC curve will be visuable

If you run all this code like confusion matrix and ROC curve will be visuable